



Valéo: Anomaly detection in Industry 4.0

Nouamane Arhachoui
Werner Dia Black
Julio Guzmán





Summary

- Introduction
- Pre-Processing
- Logistic Regression
- Neural Networks
- Ensemble Learning
- Unsupervised Learning
- Conclusion



Introduction


Valeo & the problem at hand

Introduction

Machine Learning applied to the Industry

- Production Line
- 15 stations
- One exit □ a motor

The Issue at Hand





Introduction

The Dataset

- 13 features/engine
- One output
- 30000+ entries

Many features

All features are related to data measured on a specific automaton.

One Goal

- Input analysis
 - Features
 - Correlation
- To achieve a predictive Model:

$$F(input) = output$$

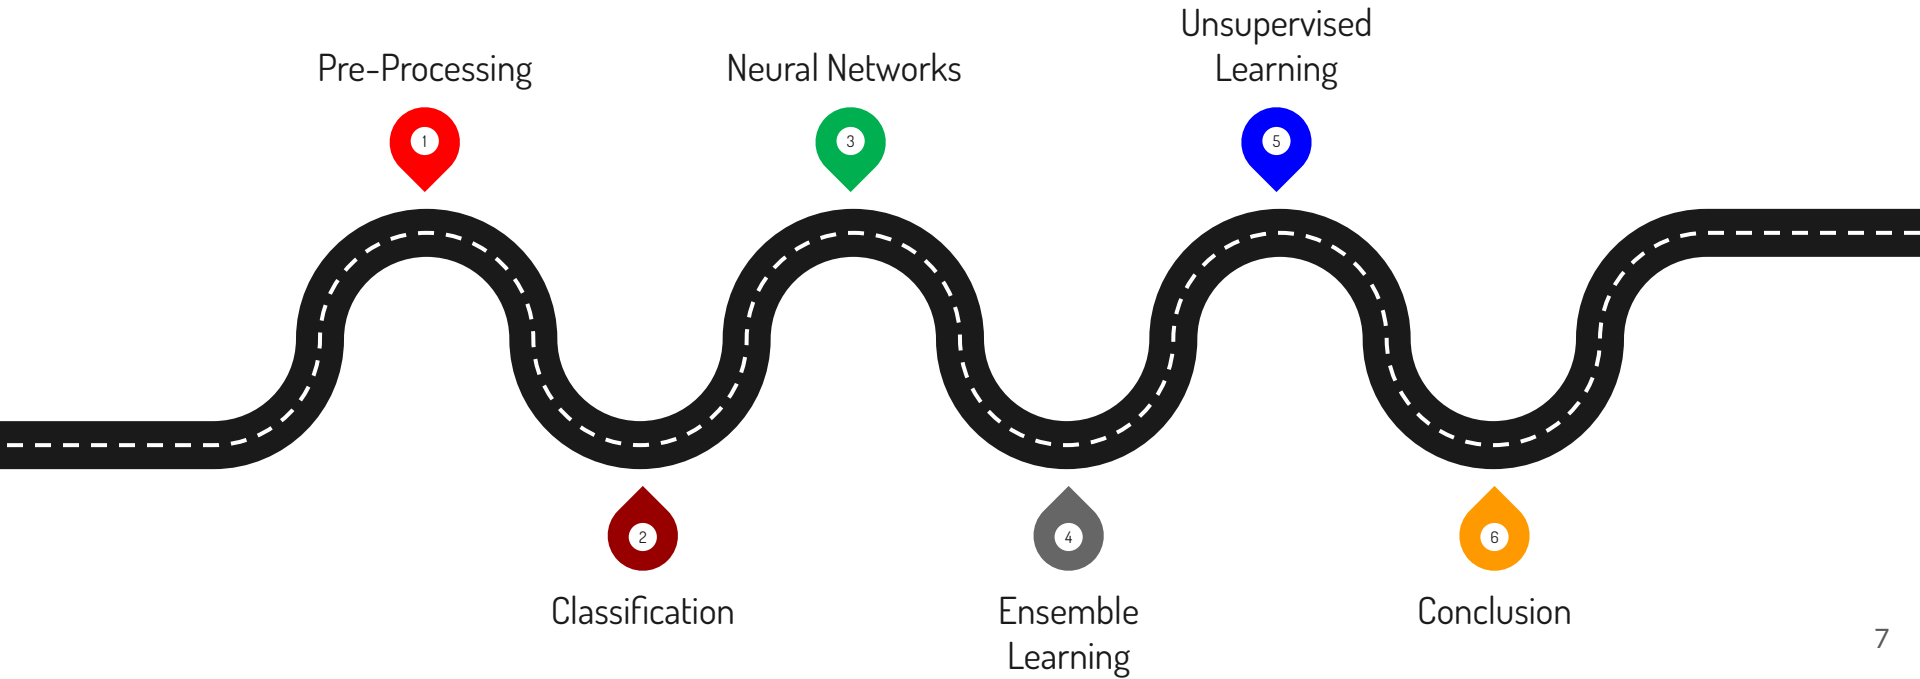
Introduction

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	PROC_TRACEINFO	34515 non-null	object
1	OP070_V_1_angle_value	34515 non-null	float64
2	OP090_SnapRingPeakForce_value	34515 non-null	float64
3	OP070_V_2_angle_value	34515 non-null	float64
4	OP120_Rodage_I_mesure_value	34515 non-null	float64
5	OP090_SnapRingFinalStroke_value	34515 non-null	float64
6	OP110_Vissage_M8_torque_value	34515 non-null	float64
7	OP100_Capuchon_insertion_mesure	15888 non-null	float64
8	OP120_Rodage_U_mesure_value	34515 non-null	float64
9	OP070_V_1_torque_value	34515 non-null	float64
10	OP090_StartLinePeakForce_value	34515 non-null	float64
11	OP110_Vissage_M8_angle_value	34515 non-null	float64
12	OP090_SnapRingMidPointForce_val	34515 non-null	float64
13	OP070_V_2_torque_value	34515 non-null	float64
14	results	34515 non-null	int64



Introduction





Pre-Processing

Adapting our Dataset to our needs

Pre-Processing

First Look

We do not take into account the column:

OP100_Capuchon_insertion_mesure with the missing values (53% is too many missing values)

Other observation

36% of the failing engines have missing values

⇒ keep the column by putting binary values in it (is the value missing or not)

```
df.columns[df.isnull().any()].tolist()
#Output : ['OP100_Capuchon_insertion_mesure']

print("Percentage of missing values = ", round(df['OP100_Capuchon_insertion_mesure'].isna().sum()*100/34514, 2), "%")
#Output: Percentage of missing values = 53.97 %

df_cleaned = df.drop(columns= ['OP100_Capuchon_insertion_mesure'])
```

```
missing_capuchon = df['OP100_Capuchon_insertion_mesure'].isna()
missing_capuchon = missing_capuchon.astype(np.int)
df['OP100_Capuchon_insertion_mesure'] = missing_capuchon
```

Pre-Processing




Statistical analysis

- Compute various statistical values for the features
- Hard to evaluate since the meaning of the dataset is unknown

Correlation

- Pearson's Correlation number gives us few correlation between parameters



Pre-Processing

Link between features

Without “capuchon_mesure”

```
values_corr = feature_engineering.get_most_correlated(df_cleaned)
values_corr[:20]
```



OP070_V_1_torque_value	OP070_V_2_torque_value	0.897
OP090_StartLinePeakForce_value	OP090_SnapRingMidPointForce_val	0.621
	OP070_V_2_torque_value	0.562
OP070_V_1_torque_value	OP090_StartLinePeakForce_value	0.543
OP090_SnapRingPeakForce_value	OP070_V_1_torque_value	0.490
	OP070_V_2_torque_value	0.482
OP110_Vissage_M8_torque_value	OP110_Vissage_M8_angle_value	-0.446
OP090_SnapRingFinalStroke_value	OP070_V_2_torque_value	-0.408
	OP090_StartLinePeakForce_value	-0.381
	OP070_V_1_torque_value	-0.381
OP090_SnapRingMidPointForce_val	OP070_V_2_torque_value	0.347
OP090_SnapRingPeakForce_value	OP090_SnapRingMidPointForce_val	0.345
	OP090_StartLinePeakForce_value	0.337
OP070_V_1_torque_value	OP090_SnapRingMidPointForce_val	0.335
OP070_V_1_angle_value	OP070_V_2_angle_value	0.187
OP120_Rodage_U_mesure_value	OP070_V_2_torque_value	0.172
	OP070_V_1_torque_value	0.169
OP090_SnapRingPeakForce_value	OP120_Rodage_U_mesure_value	0.135
	OP110_Vissage_M8_torque_value	-0.135
OP120_Rodage_I_mesure_value	OP120_Rodage_U_mesure_value	-0.119

With “capuchon_mesure” as binary

OP070_V_1_torque_value	OP070_V_2_torque_value	0.897
OP100_Capuchon_insertion_mesure	OP070_V_2_torque_value	0.668
	OP070_V_1_torque_value	0.665
OP090_StartLinePeakForce_value	OP090_SnapRingMidPointForce_val	0.621
OP090_SnapRingFinalStroke_value_exp	OP090_StartLinePeakForce_value	-0.593
OP090_StartLinePeakForce_value	OP070_V_2_torque_value	0.562
OP090_SnapRingFinalStroke_value_exp	OP070_V_2_torque_value	-0.559
OP070_V_1_torque_value	OP090_StartLinePeakForce_value	0.543
OP090_SnapRingFinalStroke_value_exp	OP070_V_1_torque_value	-0.523
OP090_SnapRingPeakForce_value	OP100_Capuchon_insertion_mesure	0.519
	OP070_V_1_torque_value	0.490
	OP070_V_2_torque_value	0.482
OP110_Vissage_M8_torque_value	OP110_Vissage_M8_angle_value	-0.446
OP100_Capuchon_insertion_mesure	OP090_StartLinePeakForce_value	0.350
OP090_SnapRingMidPointForce_val	OP070_V_2_torque_value	0.347
OP090_SnapRingPeakForce_value	OP090_SnapRingMidPointForce_val	0.345
	OP090_StartLinePeakForce_value	0.337
OP070_V_1_torque_value	OP090_SnapRingMidPointForce_val	0.335
OP100_Capuchon_insertion_mesure	OP090_SnapRingMidPointForce_val	0.250
OP090_SnapRingFinalStroke_value_exp	OP090_SnapRingMidPointForce_val	-0.223



Logistic Regression

Logistic Regression

Binary Classification

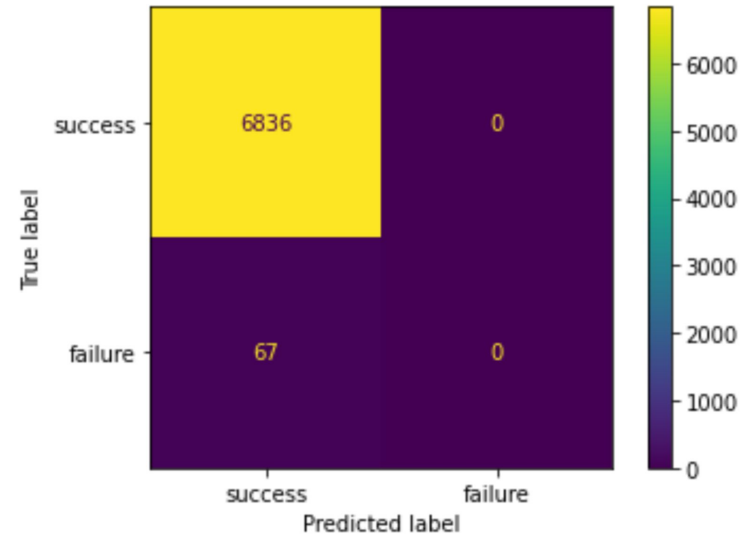
We only have two *labels* :

Pass or Not.


First, we tried Logistic Regression without regularization.

⇒ everything is predicted as a success engine

The accuracy is good but the model is bad regarding our goal



Logistic Regression



Scaled dataset	yes
Balanced dataset	yes

Only 0.9% of our Dataset is failures
⇒ SMOTE in order to balance the classes

accuracy = 0.64

idea: Grid Search Cross Validation to find the best regularization parameter to improve the model

Logistic Regression




Grid Search CV

We found the best regularization weight:

$C = 0.1$

The recall on each class is a little bit better


Scaled dataset	yes
Balanced dataset	yes





Neural Networks

Neural Networks



Neural Networks



Building a Model

- Sequential Model
- Relu as activation function and SoftMax for output
- 3 types of architecture :
 - Shallow
 - Wide
 - Deep

Training the model


- Modify Learning Rate
- Reach the Highest accuracy
- Then apply the prediction

```
def trainmodel(learning_rate):  
    model = make_sequential_model(sample_size, shallow_architecture,  
                                   learning_rate=0.001)  
    nn_file = my_path + 'nn-'+str(learning_rate)+'.h5'  
    model.summary  
  
    history = train_model(model, nn_file, X_train_balanced, y_train1, seed=5,  
                           max_epochs=200)
```


Neural Networks



Shallow[8,8,8,8,2]




Deeper[8, 8, 8, 8, 8, 8, 8, 8, 8, 2]




Neural Networks

Wide [72, 72, 72, 2]

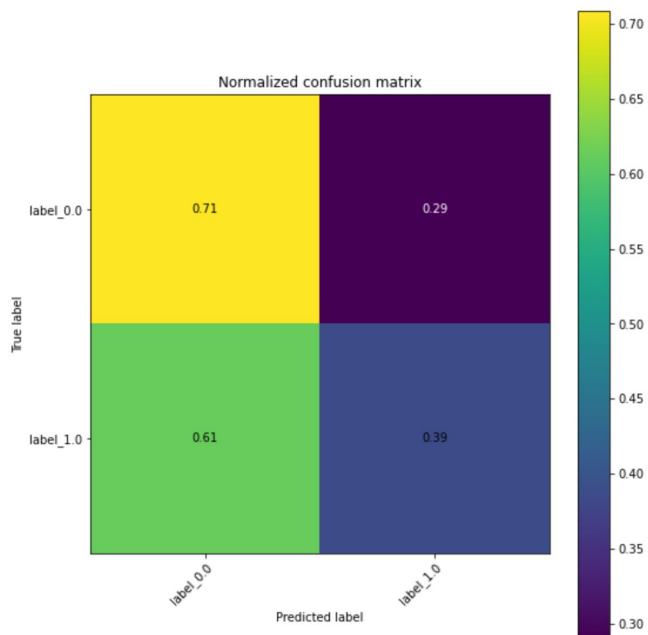


Wide and Deeper[48, 48, 48, 48, 48, 48, 48, 48, 2]




Neural Networks

For The Shallow Network [8,8,8,8,2]




For our Deeper[8,8,8,8,8,8,8,8,2]




Neural Networks



For The Wide [72, 72,72,2]



For our Hybrid Network [48,48,48,48,48,48,48,2]





Ensemble Learning

Random Forest

Balanced dataset	yes
Scaled dataset	no


```
smote = SMOTE()

X_train_balanced, y_train_balanced = smote.fit_resample(X_train, y_train)

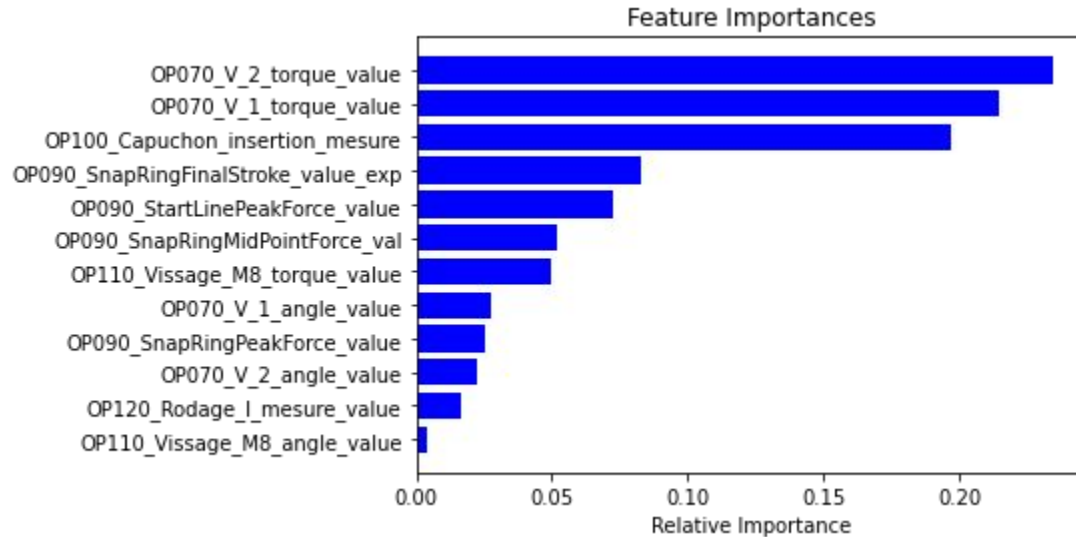
model = RandomForestClassifier(n_estimators=100,
                              criterion='gini',
                              max_leaf_nodes=16,
                              random_state=5,
                              n_jobs=-1,
                              max_features = 'auto' # sqrt(n_features)
                              )

model.fit(X_train_balanced, y_train_balanced)
```

Better results than LogisticRegression but still not satisfying



Feature importance



- Some features seem much more important than others

Randomized Search

```
forest = RandomForestClassifier(n_estimators=100, random_state = 4, n_jobs=1,)

search = RandomizedSearchCV(
    scoring = 'balanced_accuracy', # we also tried with 'f1'
    estimator=forest,
    param_distributions=param_grid,
    n_iter=50,
    verbose=2, random_state=42,
    n_jobs=-1,
    cv=5)

model = search_or_load(filename, search, X_train_balanced, y_train_balanced)
```

```
param_grid = {
    'criterion':['gini', 'entropy'],
    'max_features':list(range(1,13)),


    'max_leaf_nodes':[16, 32, 64, 128, 256, 512, 1024, 2048],
    'min_impurity_decrease' : [0, 0.001, 0.01, 0.1, 0.2],

    'max_depth':[1,10,100,1000,10000,100000],
    'min_weight_fraction_leaf' : [0.1, 0.01, 0.001, 0]
}
```

Goal:

- Trying many configurations automatically
- Get the best hyperparameters

Randomized Search



The recall on failures is lower than before

model parameters (found by RandomizedSearchCV):


```
{ 'min_weight_fraction_leaf': 0,  
  'min_impurity_decrease': 0,  
  'max_leaf_nodes': 128,  
  'max_features': 5,  
  'max_depth': 10,  
  'criterion': 'gini'}
```



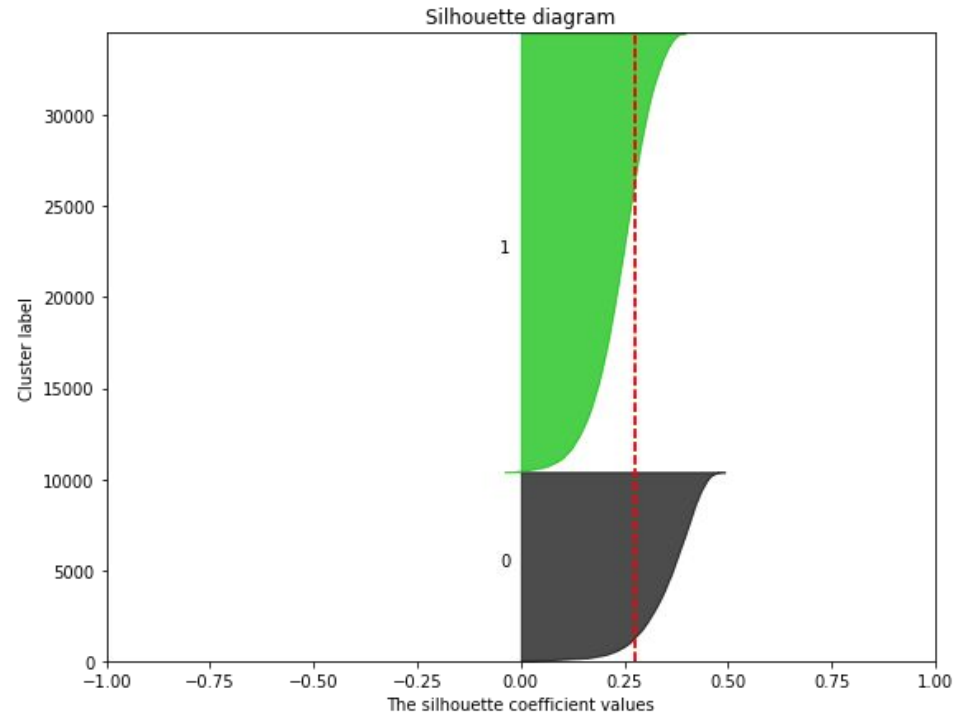
Unsupervised Learning

KMeans Clustering, Isolation Forrest and Autoencoders

K-Means Clustering





Scaled dataset	yes
Balanced dataset	no



Silhouette analysis for KMeans Clustering with clusters = 2


Avg = 0.2763

K-Means Clustering



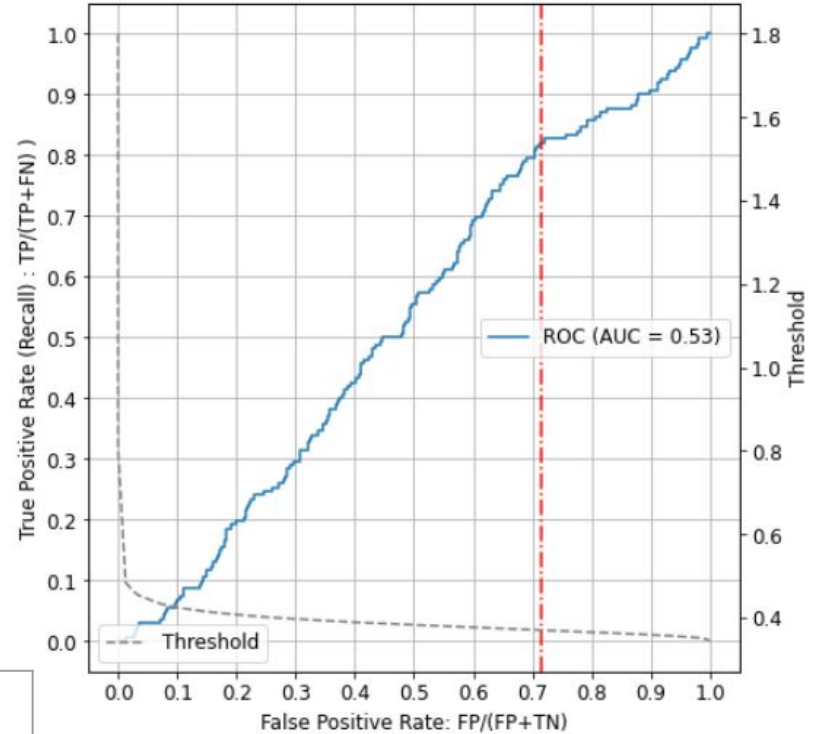
- Scaling with Standard Scaler
- K=2 and repetitions = 130
- All the dataset used to train the model

Isolation Forests



- Trees = 1500
- Threshold = 0.37

Scaled dataset	no
Balanced dataset	no

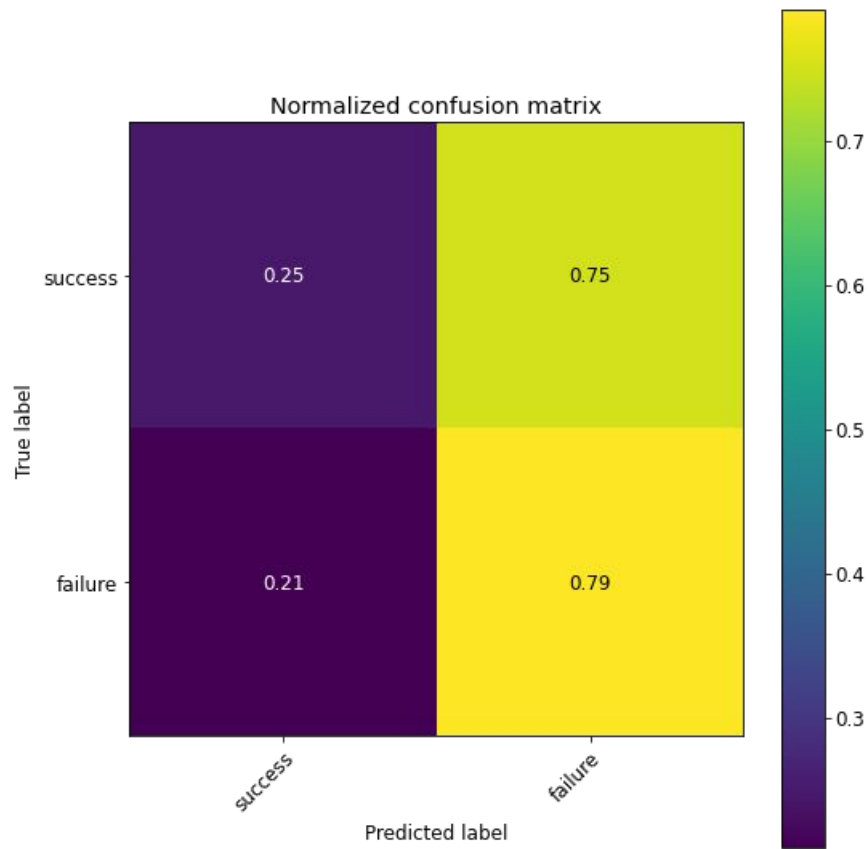


Receiver-Operating Characteristic (ROC) Curve

Isolation Forests

Scaled dataset	no
Balanced dataset	no

- Trees = 1500
- Threshold = 0.37
- Precision = 0.0087
- Recall = 0.79




Auto-Encoders

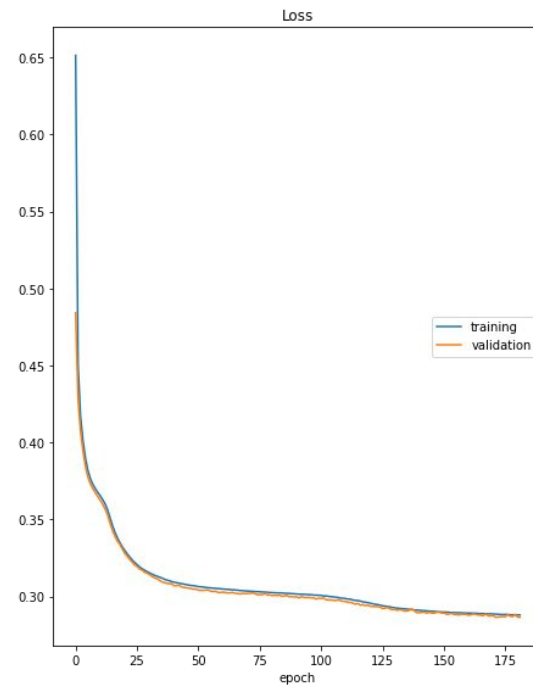


Scaled dataset	yes
Balanced dataset	no

Loss function: MSE




Learning rate = 0.01




Learning rate = 0.001

Auto-Encoders



Scaled dataset	yes
Balanced dataset	no



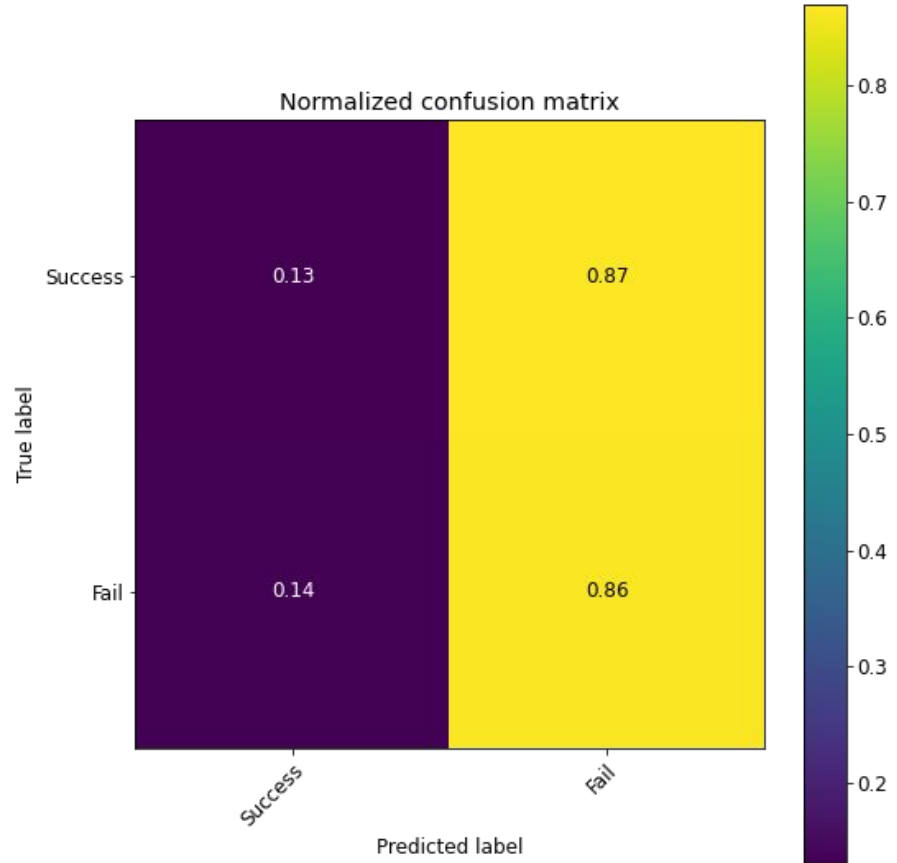
Receiver-Operating Characteristic (ROC) Curve

Threshold = 1

Auto-Encoders

Scaled dataset	yes
Balanced dataset	no

- Loss function: MSE
- Learning rate = 0.001
- Epochs = 182
- Precision = 0.0089
- Recall = 0.86





Conclusion

Comparison Between the Techniques



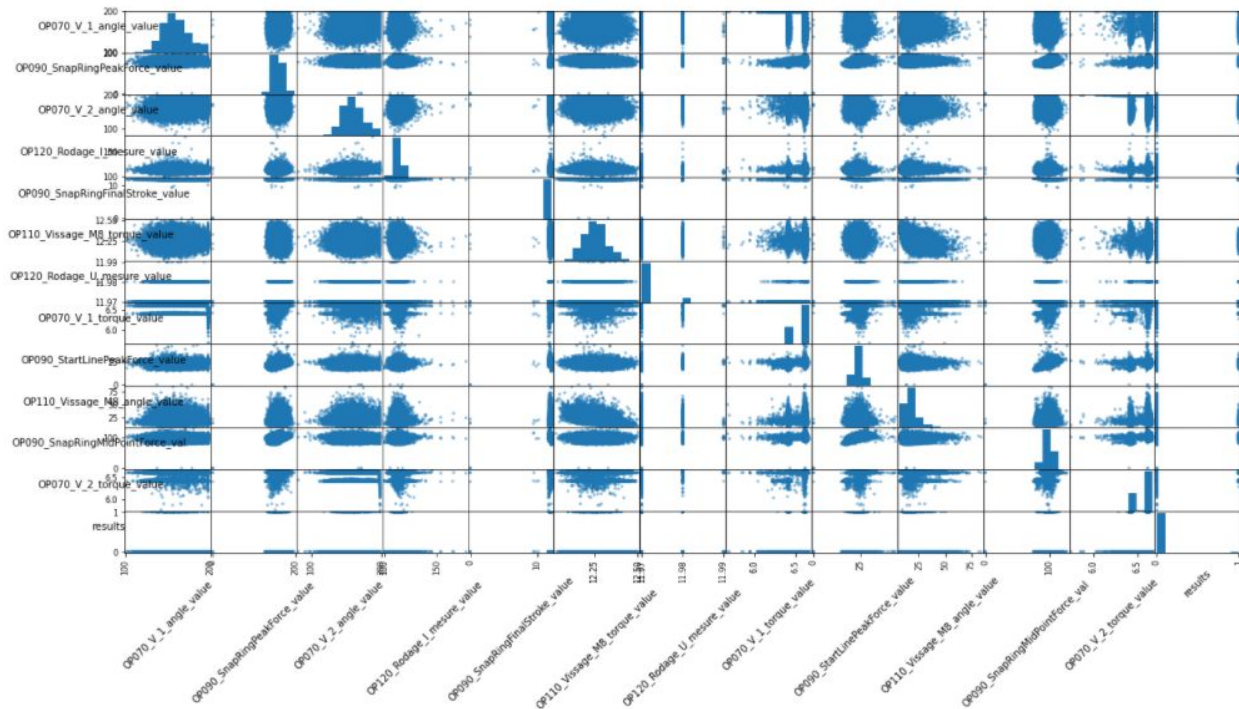
Machine Learning Technique	Precision (class: failure)	Recall (class: failure)
Logistic Regression	0.01	0.52
Neural Networks	0.01	0.30
Ensemble Learning	0.02	0.33
KMeans	0.01	0.51
Isolation Forrest	0.0087	0.79
Auto-Encoders	0.0089	0.86

THANKS!
Any questions?

Appendix


[1].Scatter matrix
obtained with the
dataset after
Pre-Processing

```
sm = scatter_matrix(df_cleaned, figsize=(20,10))  
visualization.rotate_labels(sm)  
plt.show()
```




Appendix

Dimensionality Reduction



10 largest singular values, normalized

- there is no “knee”
- hard to choose the right number of components



- all the variables are important
- maybe, try with the 10 most important

Logistic Regression

